


Long turns explain most of the LLMs' extra alignment

(a) Over a conversation (prompt P1): humans stay flat, most LLM setups converge



- Humans (p = 0.89)
- C1 all at once (p = 0.005)
- C2 turn by turn (p = 0.502)
- A3 2 agents (same model) (p = 0.001)
- A4 2 agents (2 models) (p = 0.004)

(b) Gap vs humans: raw (dashed outline) vs controlling for turn length (filled)

