


Long turns explain most of the LLMs' extra alignment

(a) Over a conversation (prompt P1): humans stay flat, most LLM setups converge



- Humans ($p = 0.89$)
- C1 all at once ($p = 0.005$)
- C2 turn by turn ($p = 0.502$)
- C3 2 agents (same model) ($p = 0.001$)
- C4 2 agents (2 models) ($p = 0.004$)

(b) Gap vs humans: raw (dashed outline) vs controlling for turn length (filled)

